



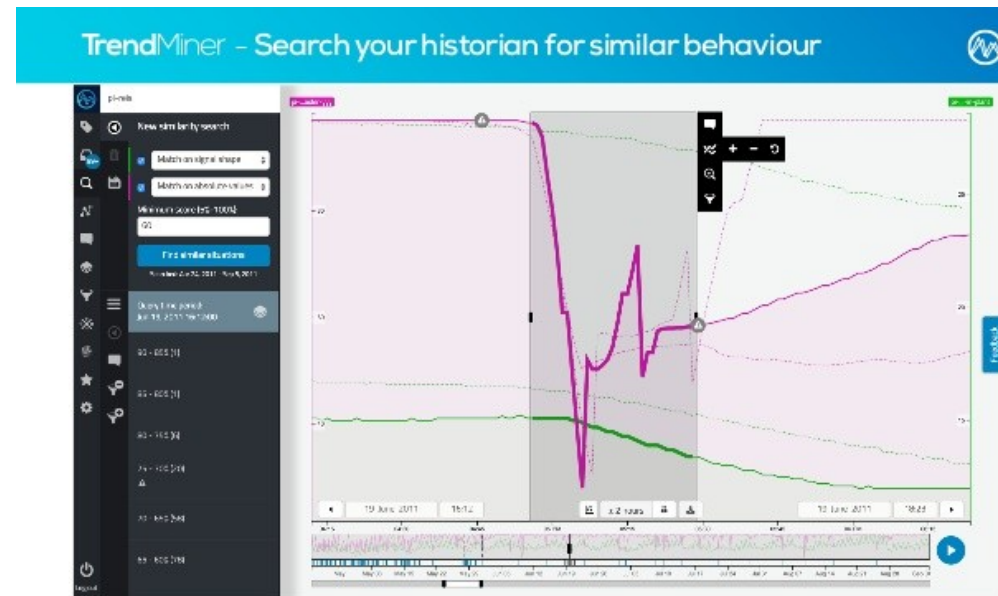
Monitoring and logging

About myself

- Studied computer science at KU Leuven
- 4 years of experience as software engineer at axxes
- Competence coach Java
- Ngdata: 10 months (Big data)
 - Hadoop
- Persgroep + truvo: 1 year
 - Spring services (micro services)
 - Handling automatisaton + dockerizing environments
- Trendminer: 2 years
 - Improve analytics for process industry
 - Timeseries data (Ignite + Cassandra)


Big Data at Trendminer

- Perform searches through timeseries data
 - Volume: Second level resolution: 400Mb/year/tag
 - Searches:
 - Finding similar patterns
 - Linear regression
 - Predicting maintenance
 - Find abnormal situations
- Goal:
 - Reduce costs
 - Improve efficiency of plants



Big Data at Telenet

- Hadoop cluster: (9Tb memory, 6 Pb storage, 30+ nodes)
- Gather + transform data to suit business
- Batch + stream processing
- Technologies
 - Cloudera stack: hdfs, hive, kudu, impala, yarn
 - Processing data: spark, kafka streams



Who are you?

- Which was the best course so far?
- What is your experience with logging?
 - Did you diagnose issues through logging?
 - How did you analyze logs? Centralized/files on the server?
 - Logging best practices?



Monitoring and logging

Overview

- Logging
 - Intro
 - What/how/when/where
 - Best practices
- Metrics
 - Application metrics
 - User metrics
 - Best practices


Logging: Intro

- Traditional way of logging
 - Simple
 - Log to files
 - Analyze: look through log file
- Challenges (functional):
 - Log the right thing
 - Too much logging is annoying
 - Too little is even worse
 - How long to keep logs
 - Storage available?
 - Regulation?


App



monolith




Infrastructure



Logging: Intro

- Microservices
 - Logs are spread over multiple servers
 - How to trace requests across services?
 - How to analyze all logs?
 - Challenges (functional):
 - Log the right thing
 - How long to keep logs
 - Challenges (technical)
 - Analyze logs efficiently
 - Trace logs across services
- => complexity increases



Logging: Intro

- Environment/context of your application is important
 - Can you (developer) access the logs
 - Direct access
 - Indirect: someone needs to send them
 - Does the application have an outbound internet connection
 - Are you allowed to see the logs
 - Regulation?
 - Sensitive information?
 - Remove unnecessary sensitive info (password/username)
 - Can you access the server?
 - Change log levels
 - Intervene when there are issues

Logging: what?

- Exceptions?
- State of the application?
- SQL logging?
- Actions performed?



Logging: what?



Logging: what?

- Exceptions?
 - Interaction with external services?
 - Unexpected behavior?
NullPointerException/ArithmeticException
 - Expected behavior?
FileNotFoundException/CancellationException
- How?
 - General exception handler (1 place) = consistency
 - Logging content
 - Exception message
 - Stacktrace
 - Context

Logging: what?

- Application state (background batch process)
 - Log in which state we are
 - history shows state transitions
 - Essential when troubleshooting
 - Reason about what went wrong
- How?
 - Log level: depending on importance
- Trade-off
 - Amount of logging
 - Enough logging to diagnose issues

Logging: what?

- SQL logging
 - Framework support (hibernate)
 - Show the actual SQL statements performed
- How?
 - At most debug level (otherwise too much spam)
 - Useful when developing/not in production
 - Often sensitive information

Logging: what?


- Actions performed?
 - Batch process/User interaction without event sourcing
 - Trace of when an operation happened
 - Be wary of amount of logging

Logging: regulations?

- We cannot log whatever we want
 - No passwords: not from the user nor the application (database access)
 - No username in the logs (depends on industry)
 - Other sensitive information
 - Impact on GDPR?
- What to do then
 - Obfuscate sensitive information
 - Make sure it is deterministic (same username to same id)
 - Think about the impact on analyzing issues (trace back to database dumps etc.)


Logging: how?

- Log files are data
 - Structure them accordingly
 - You will analyze them often
 - Using grep/awk/cat
 - Using a query engine (kibana)
- Which one do you like
 - "Could not update entry 123 as t was not found in the database"
 - "error: entry does not exist, id: 123, table: users"
 - "{type: error, data: {message: entry does not exist, id: 123, table: users}}"
- Free text is generally a bad idea
 - Try to use keywords with meaning
 - Use a different log format for stacktraces
 - Easily filter them out



Logging: how much logging?

- Logging can impact performance
- Use log levels
 - Trace/debug/info/warn/error
 - specify log levels for specific packages
 - No debug logging of frameworks
 - Might use debug logging for your app
- Can you alter log level at runtime?
 - Can use more strict log levels



Logging: reduce size of log files?

- Rolling files (date suffix):
 - Maximum size (100 Mb)
 - Maximum file history to keep (10 files)
 => predictable storage requirements
- Compress old logs (save storage)

```
niels@niels-ThinkPad-P51:~/Downloads/turay/mnt/data/trendminer-appliance-logs/docker$ ls -lh
total 1,4G
-rw-r----- 1 niels niels 164M Jul 31 22:11 docker.log
-rw-r----- 1 niels niels 38M Jul 22 02:01 docker.log-20180722
-rw-r----- 1 niels niels 2,5M Jul 23 02:01 docker.log-20180723.gz
-rw-r----- 1 niels niels 2,5M Jul 24 02:01 docker.log-20180724.gz
-rw-r----- 1 niels niels 120M Jul 25 02:01 docker.log-20180725
-rw-r----- 1 niels niels 381M Jul 26 02:02 docker.log-20180726
-rw-r----- 1 niels niels 271M Jul 27 02:02 docker.log-20180727
-rw-r----- 1 niels niels 55M Jul 28 02:01 docker.log-20180728
-rw-r----- 1 niels niels 87M Jul 29 02:01 docker.log-20180729
-rw-r----- 1 niels niels 7,3M Jul 30 02:02 docker.log-20180730.gz
-rw-r----- 1 niels niels 163M Jul 31 02:02 docker.log-20180731
-rw-rw-r-- 1 niels niels 139M Aug 27 15:54 hps.log
niels@niels-ThinkPad-P51:~/Downloads/turay/mnt/data/trendminer-appliance-logs/docker$ gunzip docker.l
og-20180730.gz
niels@niels-ThinkPad-P51:~/Downloads/turay/mnt/data/trendminer-appliance-logs/docker$ ls -lh
total 1,6G
-rw-r----- 1 niels niels 164M Jul 31 22:11 docker.log
-rw-r----- 1 niels niels 38M Jul 22 02:01 docker.log-20180722
-rw-r----- 1 niels niels 2,5M Jul 23 02:01 docker.log-20180723.gz
-rw-r----- 1 niels niels 2,5M Jul 24 02:01 docker.log-20180724.gz
-rw-r----- 1 niels niels 120M Jul 25 02:01 docker.log-20180725
-rw-r----- 1 niels niels 381M Jul 26 02:02 docker.log-20180726
-rw-r----- 1 niels niels 271M Jul 27 02:02 docker.log-20180727
-rw-r----- 1 niels niels 55M Jul 28 02:01 docker.log-20180728
-rw-r----- 1 niels niels 87M Jul 29 02:01 docker.log-20180729
-rw-r----- 1 niels niels 124M Jul 30 02:02 docker.log-20180730
-rw-r----- 1 niels niels 163M Jul 31 02:02 docker.log-20180731
-rw-rw-r-- 1 niels niels 139M Aug 27 15:54 hps.log
```

```
TimeBasedRollingPolicy">
fileNamePattern>
```

Logging: logging libraries?

- Slf4j (API)
- API implementations
 - Logback
 - Add logback.xml in resources with logging configuration
 - Log4j2/log4net
 - Commons-logging (do not use)
 - Util logging (issues when running app as executable jar)
- If you use spring
 - `<include resource="org/springframework/boot/logging/logback/base.xml"/>`
 - Add custom configuration

Logging: logback configuration?

```
<configuration>
  <include resource="org/springframework/boot/logging/logback/defaults.xml" />
  <property name="CONSOLE_LOG_PATTERN" value="${CONSOLE_LOG_PATTERN:-%clr(%d{${LOG_DATEFORMAT_PATTERN:-yyyy-MM-dd HH:mm:ss.SSS}}){faint} %clr(${LOG_LEVEL_PATTERN:-%5p}) %clr(${PID:- }){magenta} %clr(---){faint} %clr([%15.15t]){faint} %clr(%-40.40logger{39}){cyan} %clr(:){faint} %m%n${LOG_EXCEPTION_CONVERSION_WORD:-%wEx}}"/>
  <property name="FILE_LOG_PATTERN"
    value="${FILE_LOG_PATTERN:-%d{${LOG_DATEFORMAT_PATTERN:-yyyy-MM-dd HH:mm:ss.SSS}} ${LOG_LEVEL_PATTERN:-%5p} ${PID:- } --- [%t] %-40.40logger{39} : %m%n${LOG_EXCEPTION_CONVERSION_WORD:-%wEx}}"/>
  <property name="LOG_FILE" value="${LOG_FILE:-${LOG_PATH:-${LOG_TEMP:-${java.io.tmpdir:-/tmp}}}/spring.log}"/>
  <include resource="org/springframework/boot/logging/logback/console-appender.xml" />
  <include resource="org/springframework/boot/logging/logback/file-appender.xml" />
  <root level="INFO">
    <appender-ref ref="CONSOLE" />
    <appender-ref ref="FILE" />
  </root>


  <root level="WARN"/>
  <logger name="org.springframework.web" level="WARN"/>
</configuration>
```

Logging: best practices?

- Use logging framework: work against API
 - Switch implementations when necessary
- Assign correct log-level + define log levels per package
- Invest in infrastructure to query logs
 - Structure your logs
 - Use ELK stack (see further)
- Use standard log structure:
 - thread name/time/context/event/result
- Log as much as you can (I always regret too little logs)
- Create dashboards if you have access to logs


Logging: Example?

- Logging application (consultants)
 - Customize logback
 - Spring boot actuator
 - Change logging at runtime



Logging: ELK stack why?

„Can you check the errors from yesterday between 15.02 and 15.07“



Logging: ELK stack why?

- Quickly analyzing logs
 - Central place for all logs
 - Unified structure across all applications
 - Powerful query functionality
- Dashboard for overview of application status
 - Visual instead of textual information
 - Example: number of errors: go up then alert developers

Logging: ELK stack

- A standard for managing log files in microservice environment
 - Parse log files and transfer them
 - Log files are sent to central database (elasticsearch)
 - Querying/analyzing/dashboards are created in kibana

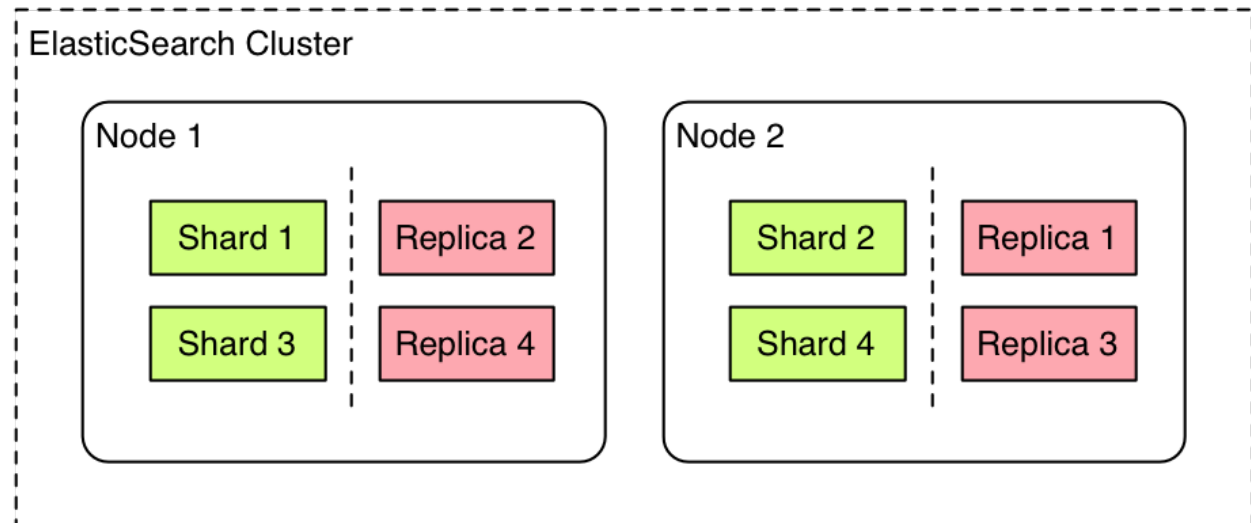


Logging: Elasticsearch

- Sharded, distributed json document store
- Supports full text search, geospatial search, ...
- Good REST API to execute searches
- Java
- Apache lucene

Logging: Elasticsearch

- Concepts:
 - Cluster: set of nodes
 - Node: instance (java process)
 - Index: collection of documents (exist many in a cluster)
 - Document: basic unit that can be searched/indexed
→ all properties in document can have a specific type
 - Shard: every index is split in multiple shards: distribute load across nodes



Demo: Elasticsearch

- Start elasticsearch:

```
$ docker run -d -p 9200:9200 -p 9300:9300 -e  
"discovery.type=single-node" elasticsearch:latest
```
- Exercises:
 - Insert a customer document with your name
 - Query the document by id
 - Query the document by name

Demo: Elasticsearch

- Insert accounts dataset
`curl -H "Content-Type: application/json" -XPOST "localhost:9200/bank/_doc/_bulk" --data-binary "@accounts.json"`
- Query documents sorted by account_number
 - alter number of results
- Query documents with mill address and only display balance
- Query documents for anyone who is 40

```
{
  "account_number": 0,
  "balance": 16623,
  "firstname": "Bradshaw",
  "lastname": "Mckenzie",
  "age": 29,
  "gender": "F",
  "address": "244 Columbus Place",
  "employer": "Euron",
  "email": "bradshawmckenzie@euron.com",
  "city": "Hobucken",
  "state": "CO"
}
```


Logging: Logstash

- Plumbing for your logs
 - Different log structure
 - Different date formats
- Many different inputs exists
 - file,kafka queue,...
- Filtering/parsing of logs
- Many different outputs exists
 - Redis, elasticsearch, file

Logging: Logstash

inputs	codecs	filters	outputs
<ul style="list-style-type: none"> • collectd • drupal_dblog • elasticsearch • eventlog • exec • file • ganglia • gelf • gemfire • generator • graphite • heroku • imap • invalid_input • irc • jmx • log4j • lumberjack • pipe • puppet_factor • rabbitmq • redis • relp • s3 • snmptrap • sqlite • sqs • stdin • stomp • syslog • tcp • twitter • udp • unix • vamishlog 	<ul style="list-style-type: none"> • cloudtrail • compress_spooler • dots • edn • edn_lines • fluent • graphite • json • json_lines • json_spooler • line • msgpack • multiline • netflow • noop • oldlogstashjson • plain • rubydebug • spool 	<ul style="list-style-type: none"> • advisor • alter • anonymize • checksum • cidr • cipher • clone • collate • csv • date • dns • drop • elapsed • elasticsearch • environment • extractnumbers • fingerprint • gelfify • geoip • grep • grok • grokdiscovery • i18n • json • json_encode • kv • metaevent • metrics • multiline • mutate • noop • prune • punct • railsparallelrequest • range 	<ul style="list-style-type: none"> • boundary • circonus • cloudwatch • csv • datadog • datadog_metrics • elasticsearch • elasticsearch_http • elasticsearch_river • email • exec • file • ganglia • gelf • gemfire • google_bigquery • google_cloud_storage • graphite • graphstastic • hipchat • http • irc • jira • juggernaut • librato • loggly • lumberjack • metriccatcher • mongodb • nagios • nagios_nsca • null • opentsdb • pagerduty • pipe

Logging: Logstash

```
input {
  file {
    path => "/tmp/access_log"
    start_position => "beginning"
  }
}

filter {
  grok {
    match => { "message" => "%{COMBINEDAPACHELOG}" }
  }
  date {
    match => [ "timestamp" , "dd/MMM/yyyy:HH:mm:ss Z" ]
  }
}

output {
  elasticsearch { hosts => ["localhost:9200"] }
  stdout { codec => rubydebug }
}
```

Logging: Grok basics

- Combines text patterns into something that matches your log
 - `%{SYNTAX:SEMANTIC}`
 - Syntax: name of pattern that will match your text
 - Semantic: identifier you give to that text being matched

Examples


55.3.244.1 GET /index.html 15824 0.043

`%{IP:client} %{WORD:method} %{URIPATHPARAM:request} %{NUMBER:bytes} %{NUMBER:duration}`

Logging: Kibana

- Configurable dashboard
- Real-time functionality
- Slice your logs stored in elasticsearch

TM Compute - Load test results



TM Compute - Load test errors

serviceId: Descending

	Count
com.trendminer.compute.load.ScatterChart_ComputeLT.fifteenTags_twoWeeks	2
com.trendminer.compute.load.ScatterChart_ComputeLT.fifteenTags_day	2
com.trendminer.compute.load.InfluenceFactor_ComputeLT.executeGlobalSearchPITeTags2Weeks_withShiftDetection	1
com.trendminer.compute.load.InfluenceFactor_ComputeLT.executeGlobalSearchPITeTags2Weeks_noShiftDetection	1
com.trendminer.compute.load.InfluenceFactor_ComputeLT.executeGlobalSearchPITeTags2Days_withShiftDetection	1
com.trendminer.compute.load.InfluenceFactor_ComputeLT.executeGlobalSearchPITeTags2Days_noShiftDetection	2
com.trendminer.compute.load.InfluenceFactor_ComputeLT.executeGlobalSearchPITeTags12Hours_withShiftDetection	1
com.trendminer.compute.load.InfluenceFactor_ComputeLT.executeGlobalSearchPITeTags12Hours_noShiftDetection	1
com.trendminer.compute.load.Histogram_ComputeLT.fiveTags_twoWeeks_5Tags	1
com.trendminer.compute.load.Histogram_ComputeLT.createHistogram_twoWeeks_manyFilters_5Tags	1

Logging: ELK stack example

- Setup ELK stack locally with docker-compose
 - Kibana
 - Elasticsearch
 - Logstash
- Configure logstash to analyze logs of consultants application
 - Errors when key constraint violation
 - Debug entries when validation fails
 - Illustrate kibana

Logging: example metric beat

- Setup with docker-compose
 - Kibana
 - Elasticsearch
 - Metricbeat: aggregate metrics from host machine
- Look at the configuration
- Show the visualizations/charts



Application metrics

Application metrics: What?

- Measure health of each service
 - Response times
 - Load on the service
 - GC/memory usage
- Measure performance of each service
 - Interaction with other service
 - Job execution time
- Measure node status
 - Allocation of memory
 - Load on the system

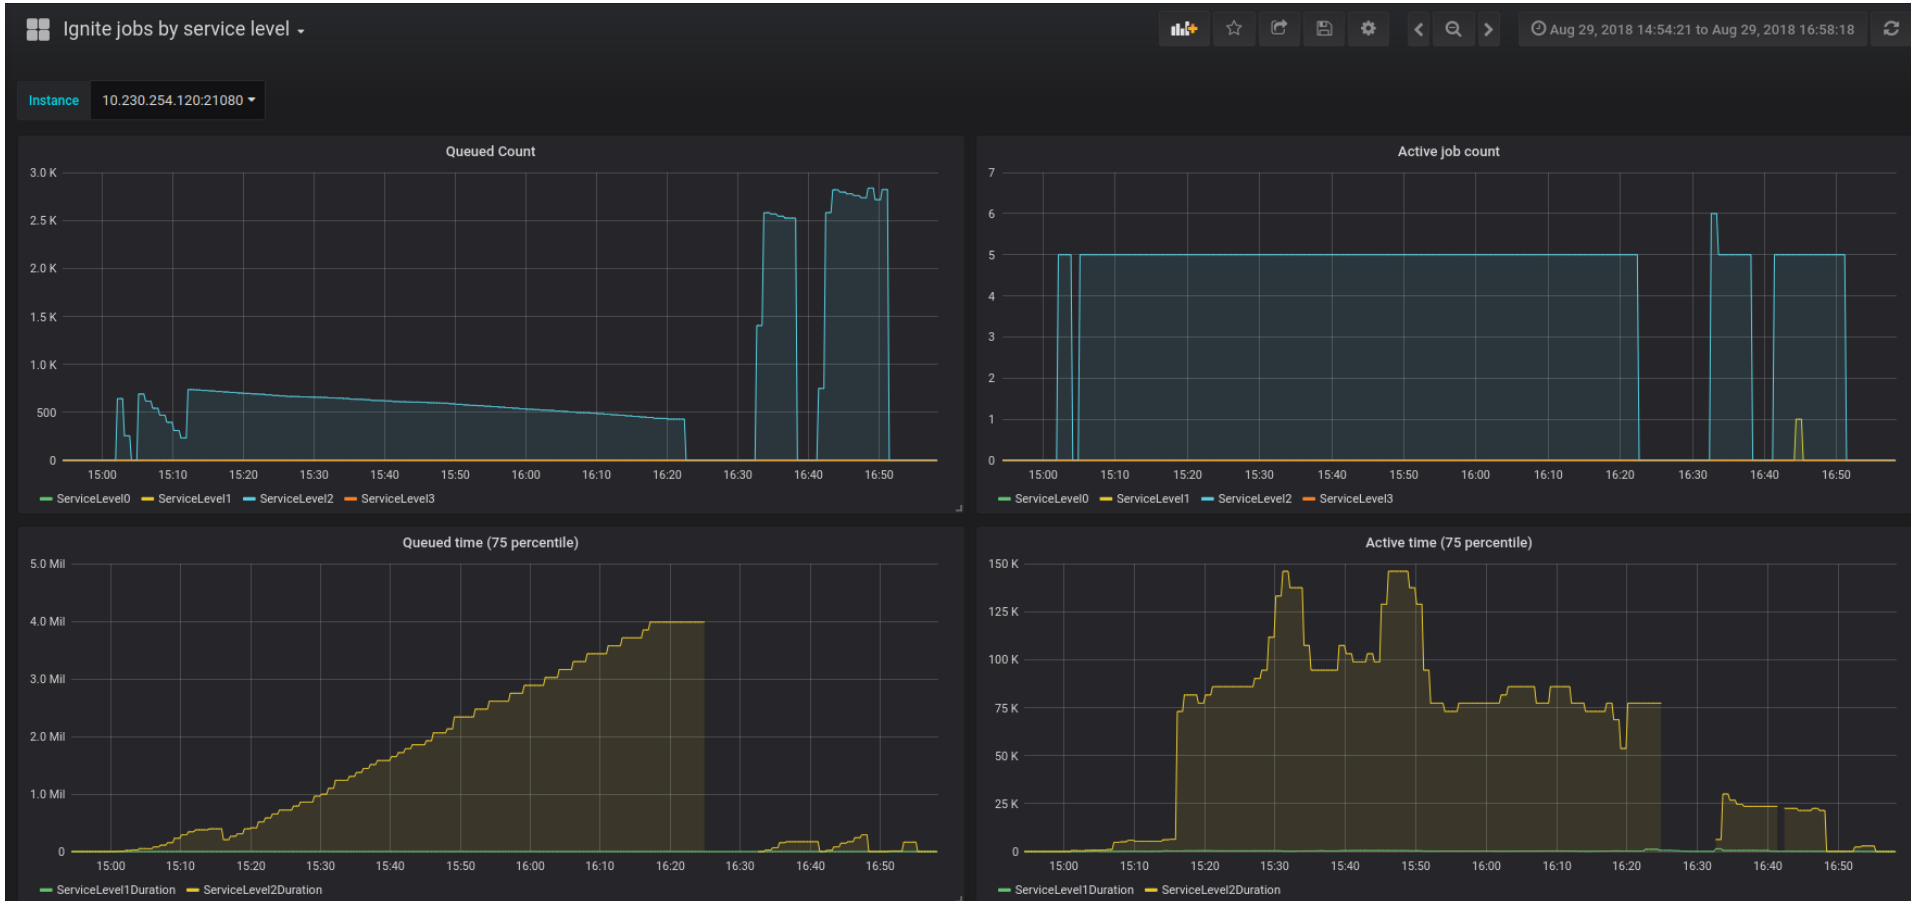
Application metrics: Why?

- Health check
 - Used to detect whether service is up (consul, eureka)
 - Load balance over multiple services
- Measure performance of application
 - detect which operations are slow and why
 - Reject/short-circuit calls if too many are queued
- Helps to diagnose why the application is slow
 - Can be that the host is overloaded (host metrics)
 - A dependent service is not responding
 - The application is doing GC (not enough memory/leak)
- Analyze non functional requirements
 - Define service levels based on them

Application metrics: How?

- Spring-boot: micrometer
 - Counters
 - Histograms: timing of operation
 - Expose these metrics
 - JMX
 - Web page
- Prometheus
 - Time series database
 - Scrape service endpoints
- Grafana
 - Visualize timeseries

Application metrics: Grafana



Application metrics: Example

- Add metrics to application
- Scrape metrics with prometheus
- Display metrics with grafana



User/Usage metrics

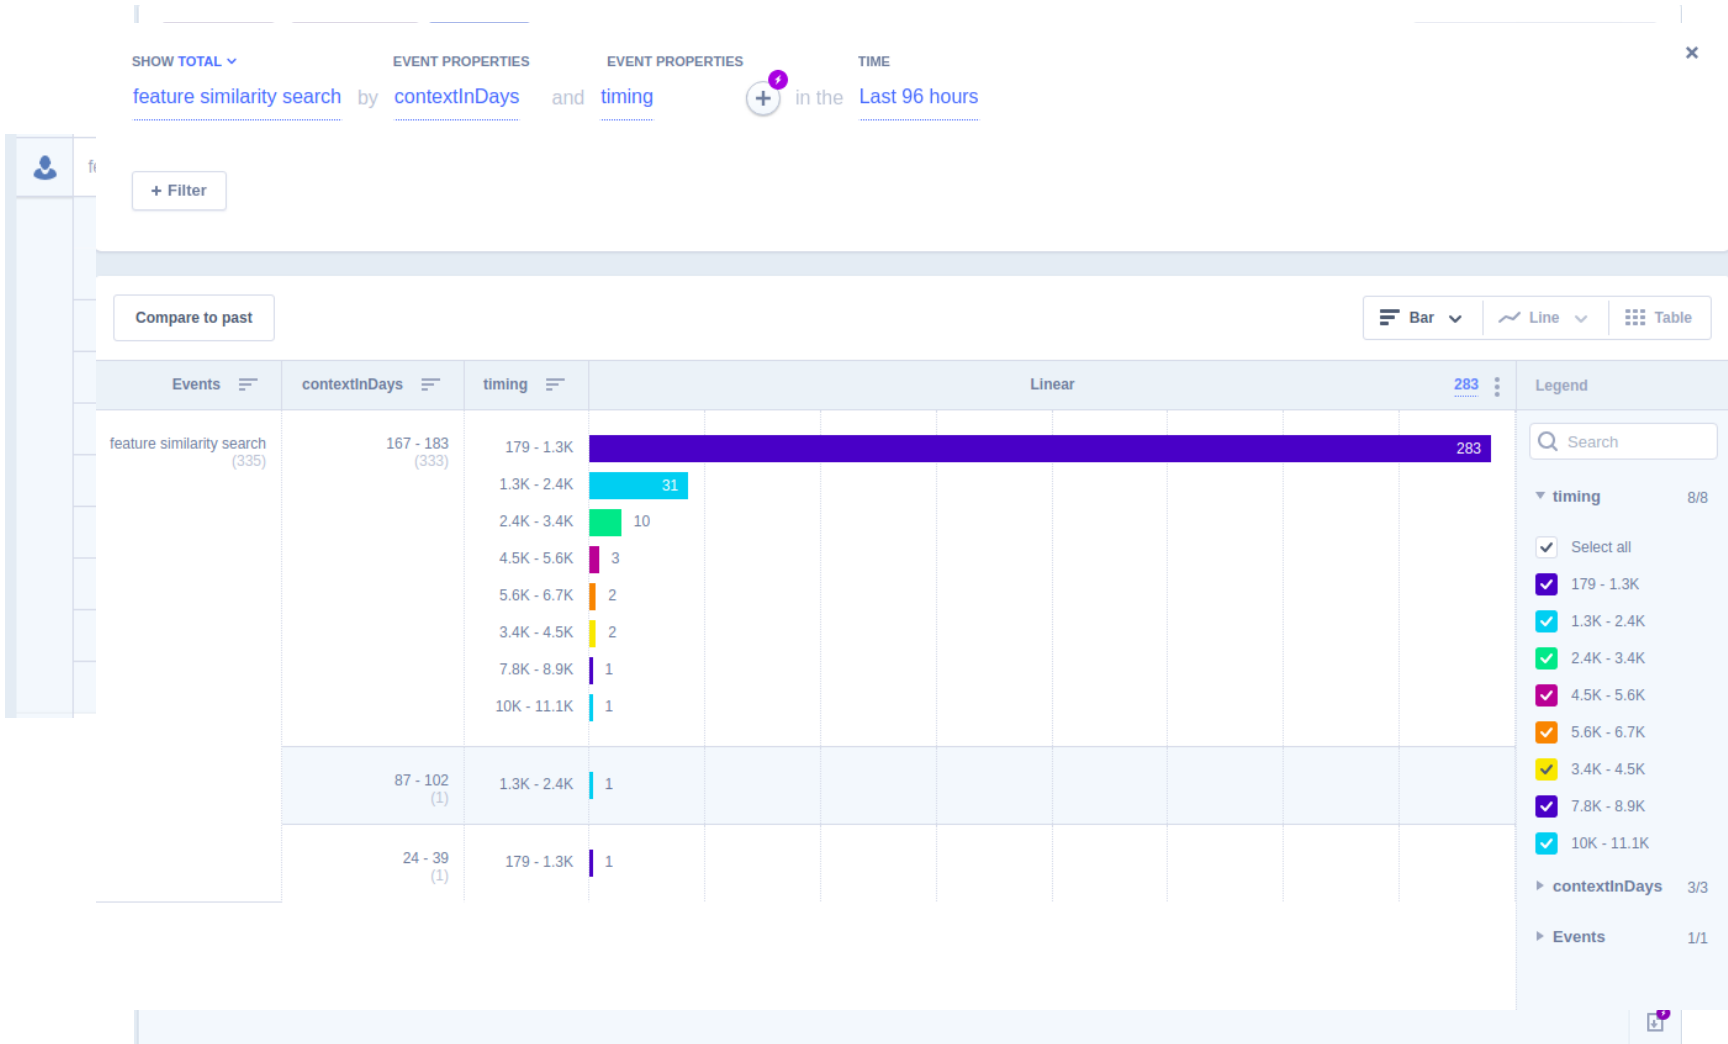
User metrics: Why?

- Gain insights into how users use the application
 - Useful when optimizing functionality
 - Optimize for 95% of use (e.g. Recommendations)
 - Analyze A/B testing: what is preferred
 - Gives info in which feature is rarely used
 - User cannot find it? → make it easily accessible
 - User does not know it? → document it
- Create user profiles
 - User who likes sports/politics/financial news
 - Based on this profile you can personalize the content

User metrics: How?

- Google analytics/mixpanel
- Log actions that a user takes
 - Provide context information: who + what + extra properties
 - Richer content than is possible in metrics
- Similar to how logging is handled
 - Central library that logs events
 - Use a pre-defined structure for all events
 - Easy analysis afterwards

User metrics: Example



**DO YOU HAVE
ANY QUESTIONS?**

**IF YOU HAVE ANY QUESTIONS IT MEANT
THAT YOU WERE NOT LISTENING SO WE ARE
NOT GOING TO ANSWER ANY QUESTIONS**

memegenerator.es

Feedback

<https://bit.ly/2P83fYU>